

**ASSIGNMENT-3**  
**ABALONE AGE**  
**PREDICTION**

Assignment Date	21 /10/2022
Student Name	PAVITHRA K
Student Roll Number	61772021T306
Maximum Marks	2 Marks

**Description:-** Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

**Task-1**

Download and Load Dataset

Download the data set:

[abalone.csv](#)

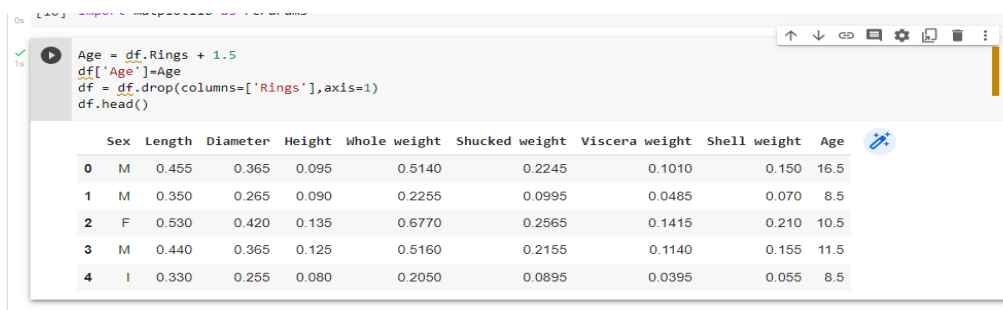
**Task-2:**

**Load the Dataset:**

Solution:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib as rcParams

df=pd.read_csv('abalone.csv')
df.head()
```

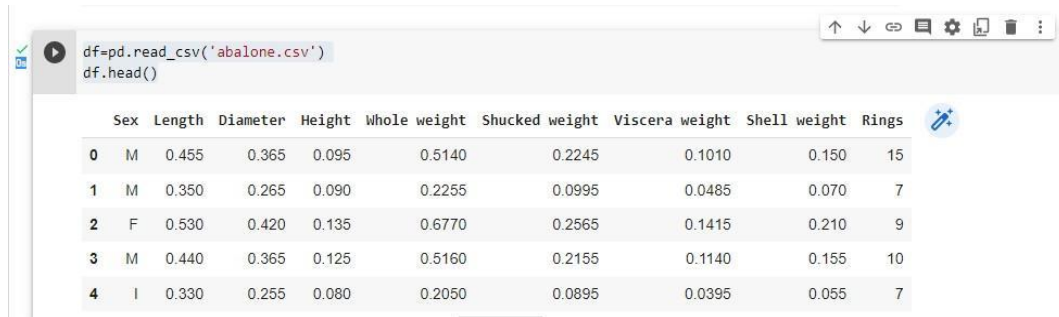


```
Age = df.Rings + 1.5
df['Age']=Age
df = df.drop(columns=["Rings"],axis=1)
df.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

### ASSIGNMENT-3 DATA VISUALIZATION AND ANALYSIS

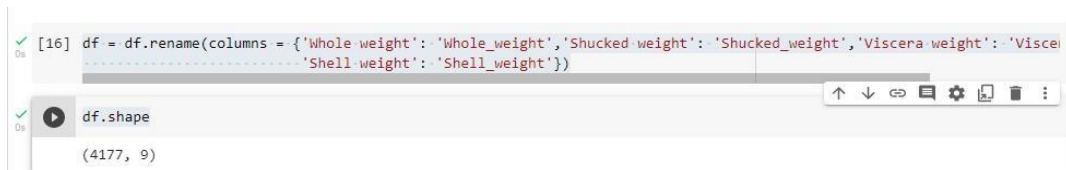
```
df=pd.read_csv('abalone.csv')  
df.head()
```



The screenshot shows a Jupyter Notebook cell with the code `df=pd.read_csv('abalone.csv')` and `df.head()`. The output is a table with 10 columns: Sex, Length, Diameter, Height, Whole weight, Shucked weight, Viscera weight, Shell weight, and Rings. The first 5 rows are displayed.

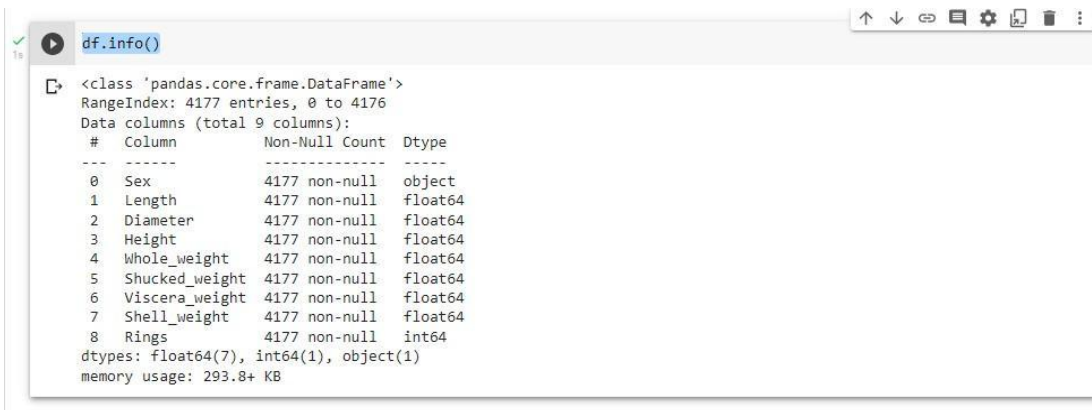
	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
df = df.rename(columns = {'Whole weight': 'Whole_weight','Shucked weight':  
'Shucked_weight','Viscera weight': 'Viscera_weight','Shell weight':  
'Shell_weight'})  
df.shape
```



The screenshot shows two Jupyter Notebook cells. The first cell contains the code to rename columns: `df = df.rename(columns = {'Whole weight': 'Whole_weight', 'Shucked weight': 'Shucked_weight', 'Viscera weight': 'Viscera_weight', 'Shell weight': 'Shell_weight'})`. The second cell contains `df.shape`, which outputs `(4177, 9)`.

```
df.info()
```



The screenshot shows a Jupyter Notebook cell with the code `df.info()`. The output provides a summary of the DataFrame, including the number of entries (4177), the number of columns (9), and the data types for each column.

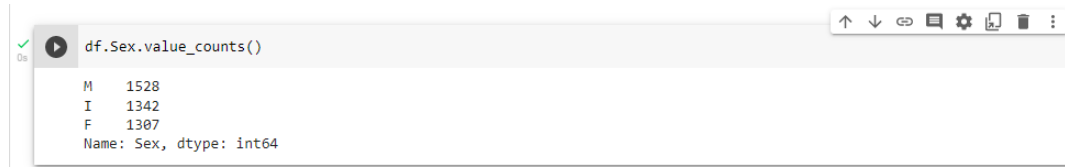
```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4177 entries, 0 to 4176  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   Sex              4177 non-null   object  
1   Length           4177 non-null   float64  
2   Diameter         4177 non-null   float64  
3   Height           4177 non-null   float64  
4   Whole_weight     4177 non-null   float64  
5   Shucked_weight   4177 non-null   float64  
6   Viscera_weight   4177 non-null   float64  
7   Shell_weight     4177 non-null   float64  
8   Rings            4177 non-null   int64  
dtypes: float64(7), int64(1), object(1)  
memory usage: 293.8+ KB
```

```
df.Sex.unique()
```



The screenshot shows a Jupyter Notebook cell with the code `df.Sex.unique()`. The output is an array of unique values: `array(['M', 'F', 'I'], dtype=object)`.

```
df.Sex.value_counts()
```



The screenshot shows a Jupyter Notebook cell with the code `df.Sex.value_counts()`. The output is a Series showing the count of each sex: `M: 1528, I: 1342, F: 1307`. The dtype is `int64`.

## ASSIGNMENT-3 ABALONE AGE PREDICTION

### Task-3:

### 3. Perform Below Visualizations.


- Univariate Analysis
- Bi - Variate Analysis
- Multi - Variate Analysis

#### Univariate Analysis:

#### SOLUTION:

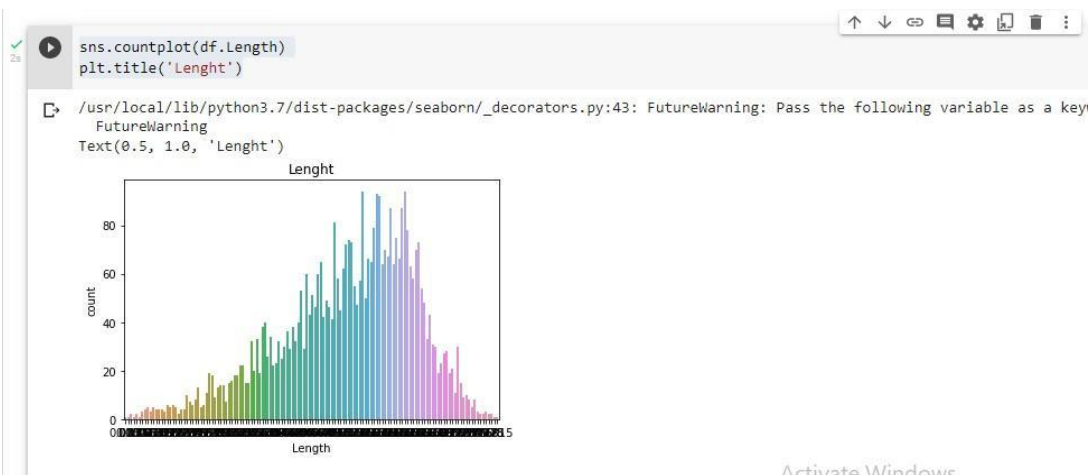
```
sns.displot(df.Age,color='Pink')
```

<seaborn.axisgrid.FacetGrid at 0x7f1974763310>



Activate Windows

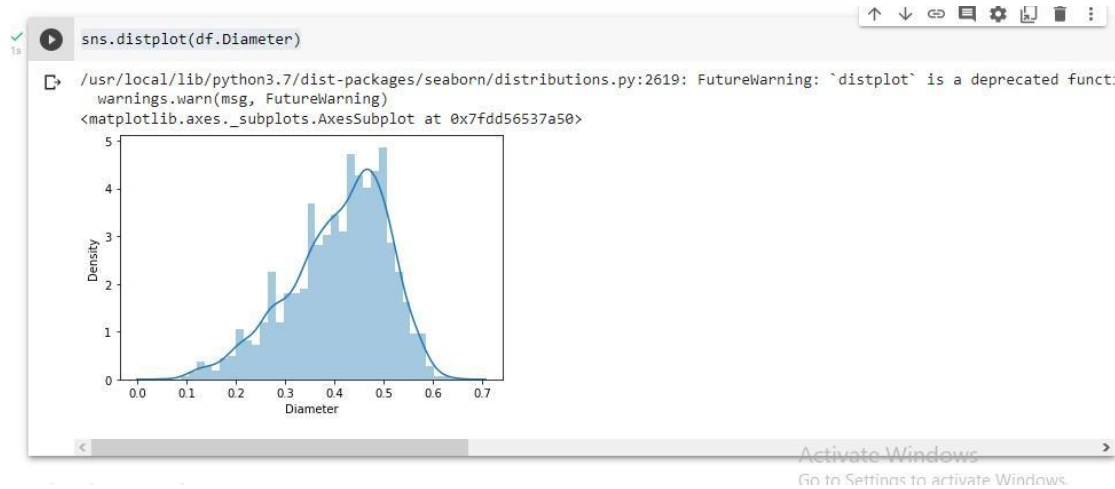
```
sns.countplot(df.Length)  
plt.title('Lenght')
```



Activate Windows

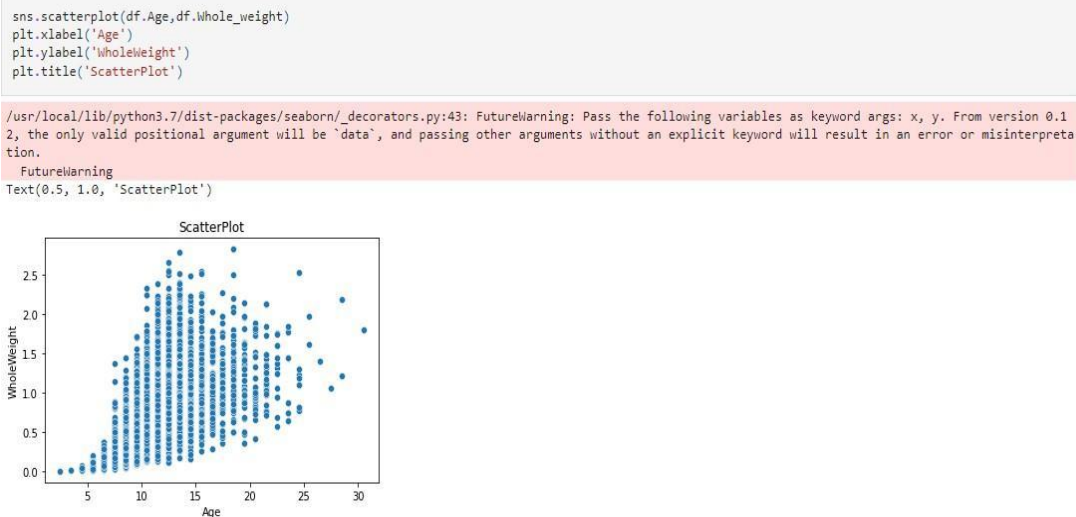
## ASSIGNMENT-3 DATA VISUALIZATION AND

```
sns.distplot(df.Diameter)
```

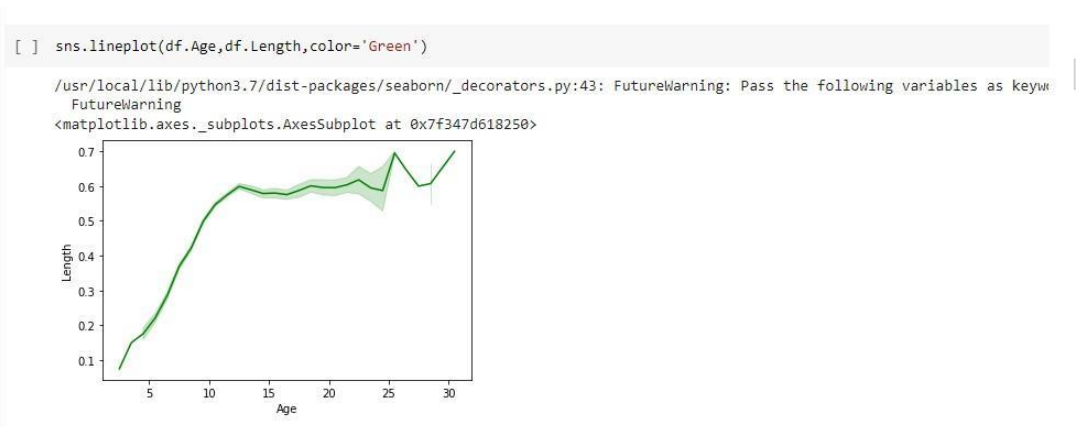


### Bi-Variate Analysis:

```
sns.scatterplot(df.Age,df.Whole_weight)  
plt.xlabel('Age')  
plt.ylabel('WholeWeight')  
plt.title('ScatterPlot')
```

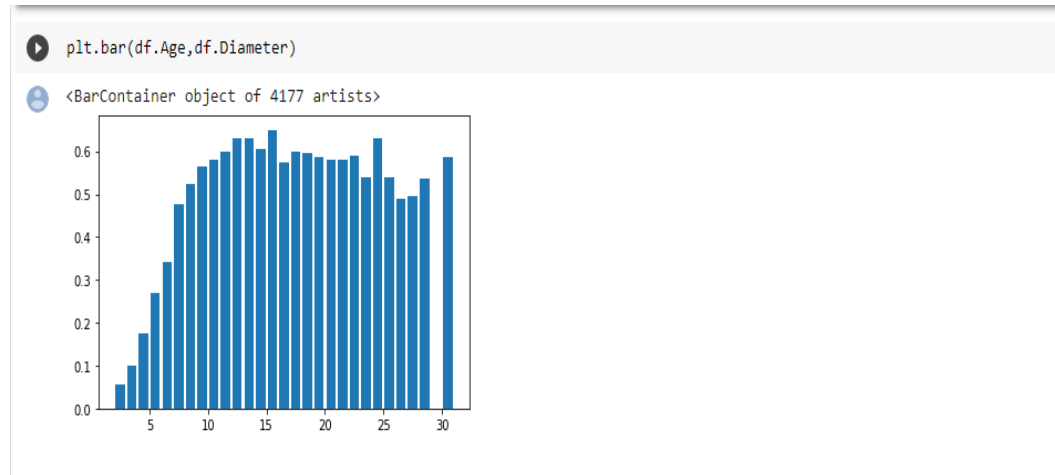


```
sns.lineplot(df.Age,df.Length,color='Green')
```



### ASSIGNMENT-3 ABALONE AGE PREDICTION

```
plt.bar(df.Age,df.Diameter)
```



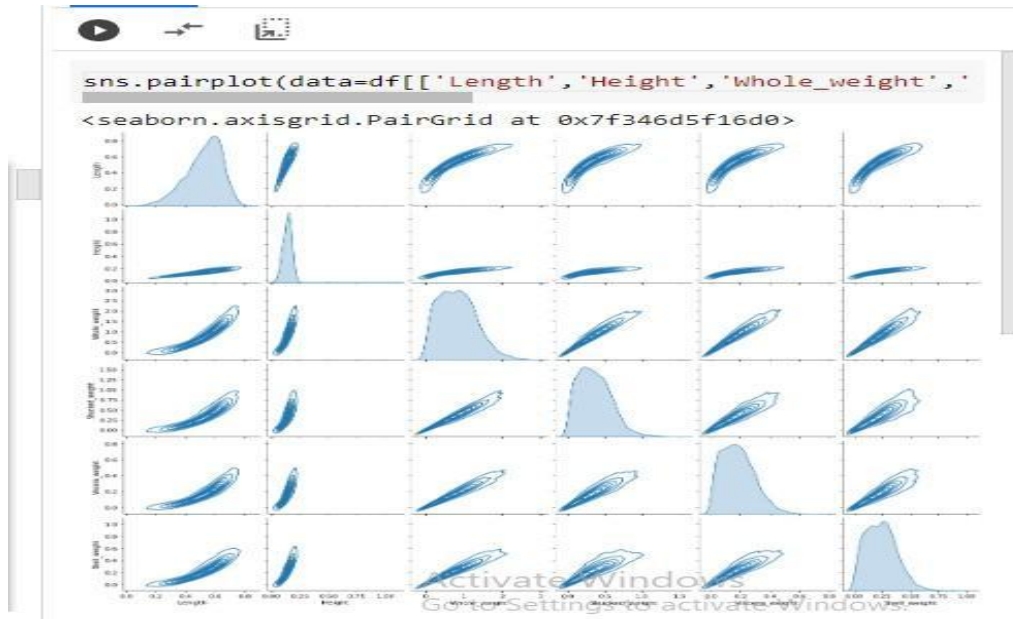
### Multi-Variate Analysis:

```
multi=df.hist(figsize=(10,10))
```

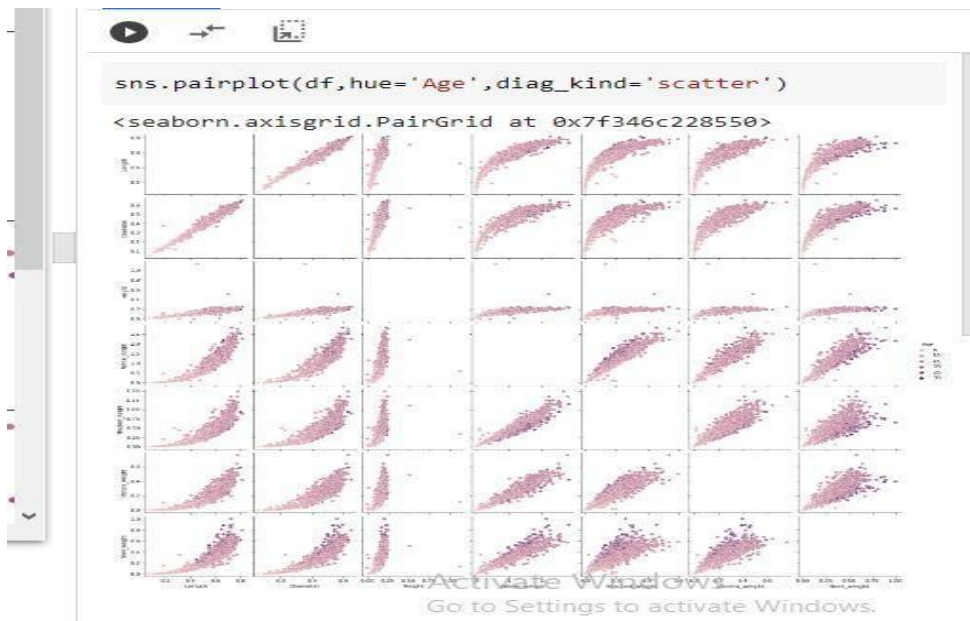


### ASSIGNMENT-3 DATA VISUALIZATION AND REGRESSION

```
sns.pairplot(data=df[['Length','Height','Whole_weight','Shucked_weight','Viscera_weight',  
'Shell_weight']],kind='kde')
```



```
sns.pairplot(df,hue='Age',diag_kind='scatter')
```



## ASSIGNMENT-3 DATA VISUALIZATION AND PREPROCESSING

### Descriptive statistics

`df.describe()`

```
df.describe()
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	11.433684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	2.500000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	9.500000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	10.500000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	12.500000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	30.500000

### Handle missing data

`df.isnull().any()` #no missing data

```
df.isnull().any()
```

Sex	False
Length	False
Diameter	False
Height	False
Whole_weight	False
Shucked_weight	False
Viscera_weight	False
Shell_weight	False
Age	False
dtype:	bool

### Outliers Replacement

`sns.boxplot(df.Age)`



```

In [ ]: q1=df.Age.quantile(0.25)
        q3=df.Age.quantile(0.75)

In [ ]: IQR = q3-q1

In [ ]: upper_limit=q3 + 1.5 * IQR
        lower_limit=q1 - 1.5 * IQR

In [ ]: upper_limit,lower_limit

Out[ ]: (17.0, 5.0)

In [ ]: df.Age.median()

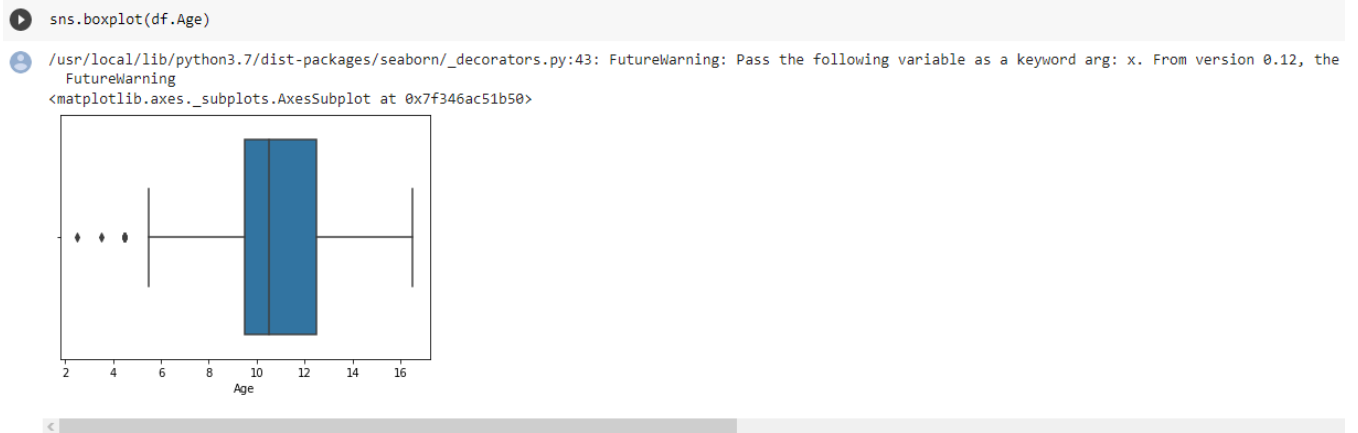
Out[ ]: 10.5

In [ ]: df.Age=np.where(df.Age>upper_limit,10.5,df.Age) #Median=10.5

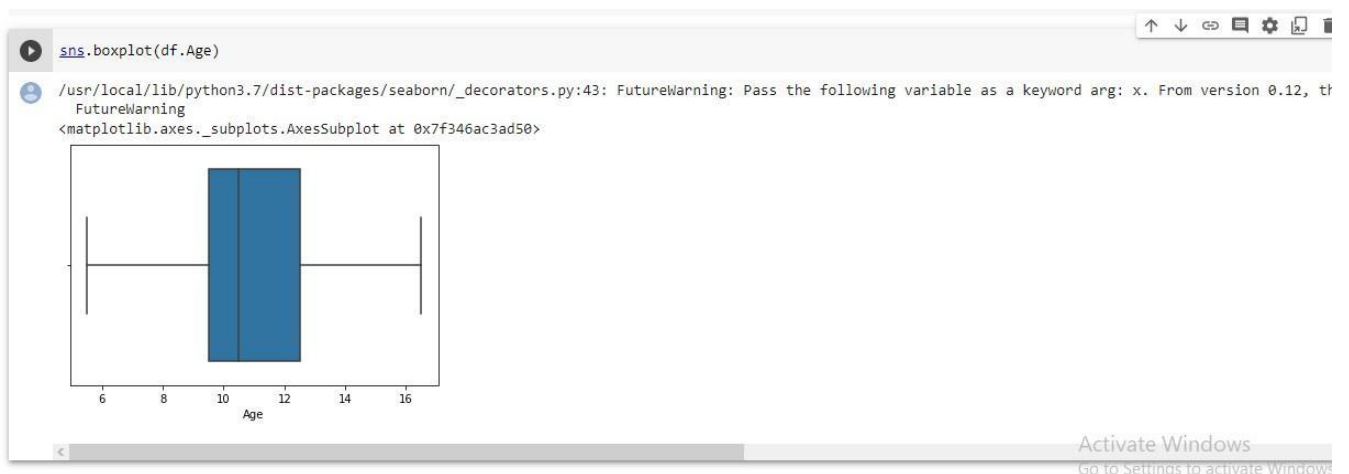
In [ ]: sns.boxplot(df.Age)

```

sns.boxplot(df.Age)



df.Age=np.where(df.Age<lower\_limit,10.5,df.Age) #Median=10.5





### ASSIGNMENT-3 DATA VISUALIZATION AND REGRESSION

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()  
df.Sex=le.fit_transform(df.Sex)  
df.head()
```

df.head()

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

### Split the data into dependent and independent variables

```
y=df['Age']  
y
```

```
#y - target columns  
#X - predicting columns  
[ ] y=df['Age']  
y  
0      16.5  
1       8.5  
2      10.5  
3      11.5  
4       8.5  
...  
4172    12.5  
4173    11.5  
4174    10.5  
4175    11.5  
4176    13.5  
Name: Age, Length: 4177, dtype: float64
```

```
X=df.drop(columns=['Age'],axis=1)  
X.head()
```

X=df.drop(columns=['Age'],axis=1)  
X.head()

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055

## Scale the independent variables

```
from sklearn.preprocessing import scale
df2=pd.DataFrame(scale(X),columns=X.columns)
df2.head()
```

```
[ ] from sklearn.preprocessing import scale
```

```
df2=pd.DataFrame(scale(X),columns=X.columns)
df2.head()
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
0	1.151980	-0.574558	-0.432149	-1.064424	-0.641898	-0.607685	-0.726212	-0.638217
1	1.151980	-1.448986	-1.439929	-1.183978	-1.230277	-1.170910	-1.205221	-1.212987
2	-1.280690	0.050033	0.122130	-0.107991	-0.309469	-0.463500	-0.356690	-0.207139
3	1.151980	-0.699476	-0.432149	-0.347099	-0.637819	-0.648238	-0.607600	-0.602294
4	-0.064355	-1.615544	-1.540707	-1.423087	-1.272086	-1.215968	-1.287337	-1.320757

## Split the data data into training and testing

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(df2,y,test_size=0.3,random_state=1)
X_train.shape,X_test.shape
```

```
y_train.shape,y_test.shape
```

```
((2923,), (1254,))
```

## Build the model

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression() #Linear Regression Model
from sklearn.linear_model import Ridge
r=Ridge() #Ridge Regression Model
from sklearn.linear_model import Lasso
l=Lasso() #Lasso Regression Model
```

```
[2] from sklearn.linear_model import LinearRegression
lr=LinearRegression() #Linear Regression Model
```

```
[4] from sklearn.linear_model import Ridge
r=Ridge() #Ridge Regression Model
```

```
[1] from sklearn.linear_model import Lasso
l=Lasso() #Lasso Regression Model
```

## Train the model

```
lr.fit(X_train,y_train) #Training lr model
pred1_train=lr.predict(X_train)
pred1_train
r.fit(X_train,y_train) #Training r model
pred2_train=r.predict(X_train)
pred2_train
l.fit(X_train,y_train) #Training l model
```

```
[ ] lr.fit(X_train,y_train) #Training lr model
LinearRegression()

[ ] pred1_train=lr.predict(X_train)
pred1_train

array([11.37532295, 10.8623978 , 10.98473747, ...,  8.47235413,
       10.0771839 ,  8.2997195 ])
```

```
[ ] r.fit(X_train,y_train) #Training r model
Ridge()

[ ] pred2_train=r.predict(X_train)
pred2_train

array([11.37024121, 10.86610153, 10.98923414, ...,  8.47158189,
       10.08029538,  8.29939996])
```

```
[ ] l.fit(X_train,y_train) #Training l model
Lasso()
```

Activate Windows  
Go to Settings to activate Windows

```
pred3_train=l.predict(X_train)
pred3_train
```

```
pred3_train=l.predict(X_train)
pred3_train

array([10.90661081, 10.94589013, 10.96552979, ..., 10.19958302,
       10.86733149, 10.14066404])
```

↑ ↓ ↻ 🗨 ✎ 🖨

## Test the model

y\_test

```
y_test

17      11.5
1131     9.5
299     10.5
1338    11.5
2383    10.5
...
802     10.5
3016     8.5
2806     9.5
2500     9.5
2814     5.5
Name: Age, Length: 1254, dtype: float64
```

```

pred1=lr.predict(X_test)
pred1
pred2=r.predict(X_test)
pred2
pred3=l.predict(X_test)
pred3

```

```

▶ pred1=lr.predict(X_test)
pred1                                     #Testing data using linear regression model
array([ 9.82570208, 10.03404396,  9.28563548, ..., 10.0965599 ,
        10.66920195,  7.77742412])

[ ] pred2=r.predict(X_test)
pred2                                     #Testing data using ridge model
array([ 9.82297354, 10.04038957,  9.28565669, ..., 10.10174311,
        10.66576784,  7.77789848])

[ ] pred3=l.predict(X_test)
pred3                                     #Testing data using lasso model
array([10.59237624, 10.96552979, 10.35670031, ..., 10.69057454,
        10.92625047, 10.04246574])

```

```

age_pred = pd.DataFrame({'Actual_value':y_test,'Predicted_value_using_lr':pred1,'Predicted_value_using_r':pred2,'Predicted_value_using_l':pred3})
age_pred.head()

```

```
[ ] age_pred = pd.DataFrame({'Actual_value':y_test,'Predicted_value_using_lr':pred1,'Predicted_value_using_r':pred2,'Predicted_value_using_l':pred3})
age_pred.head()
```

	Actual_value	Predicted_value_using_lr	Predicted_value_using_r	Predicted_value_using_l
17	11.5	9.825702	9.822974	10.592376
1131	9.5	10.034044	10.040390	10.965530
299	10.5	9.285635	9.285657	10.356700
1338	11.5	11.109891	11.111671	11.044088
2383	10.5	10.901944	10.905969	10.788773

## Measure the performance using metrics

```

from sklearn import
metrics #R2-square
#Testing accuracy of linear regression, ridge, lasso
print(metrics.r2_score(y_test,pred1))
print(metrics.r2_score(y_test,pred2))
print(metrics.r2_score(y_test,pred3))

```

```

[ ] from sklearn import metrics

▶ #R2-square
#Testing accuracy of linear regression, ridge, lasso

print(metrics.r2_score(y_test,pred1))
print(metrics.r2_score(y_test,pred2))
print(metrics.r2_score(y_test,pred3))

0.4162940378151394
0.41640627795250973
0.17272068414915298

```

### ASSIGNMENT-3 DATA VISUALIZATION AND PREPROCESSING

#R2-square

#Training accuracy of linear regression, ridge, lasso

```
print(metrics.r2_score(y_train,pred1_train))
```

```
print(metrics.r2_score(y_train,pred2_train))
```

```
print(metrics.r2_score(y_train,pred3_train))
```

```
#R2-square
#Training accuracy of linear regression, ridge, lasso

print(metrics.r2_score(y_train,pred1_train))
print(metrics.r2_score(y_train,pred2_train))
print(metrics.r2_score(y_train,pred3_train))
```

```
0.40173116413670873
0.40172280022100826
0.17472314547809642
```

## MSE(Mean square error)

Testing accuracy of linear regression, ridge, lasso

```
print(metrics.mean_squared_error(y_test,pred1))
```

```
print(metrics.mean_squared_error(y_test,pred2))
```

```
print(metrics.mean_squared_error(y_test,pred3))
```

```
## MSE(Mean square error)
#Testing accuracy of linear regression, ridge, lasso

print(metrics.mean_squared_error(y_test,pred1))
print(metrics.mean_squared_error(y_test,pred2))
print(metrics.mean_squared_error(y_test,pred3))
```

```
3.066910254318059
3.0663205217291396
4.346694365552255
```

## RMSE

#Testing accuracy of linear regression, ridge, lasso

```
print(np.sqrt(metrics.mean_squared_error(y_test,pred1)))
```

```
print(np.sqrt(metrics.mean_squared_error(y_test,pred2)))
```

```
print(np.sqrt(metrics.mean_squared_error(y_test,pred3)))
```

```
## RMSE
#Testing accuracy of linear regression, ridge, lasso

print(np.sqrt(metrics.mean_squared_error(y_test,pred1)))
print(np.sqrt(metrics.mean_squared_error(y_test,pred2)))
print(np.sqrt(metrics.mean_squared_error(y_test,pred3)))
```

```
1.751259619336339
1.7510912374085879
2.084872745649541
```

